

P15. D–HAL: Distributed Hierarchical Adversarial Learning for Multi–Agent Interaction in Autonomous Intersection Management

Guangzhou Li; Jianping Wu; Yujing He 2023 arXiv preprint PDF pages: 17 Relevance: 73/100

1. 论文全景速览与核心贡献 / High-Level Overview & Core Contributions

研究背景与痛点 / Background & Pain Points

这篇论文位于“Distributed hierarchical learning for AIM”方向。对你的 JSSP-based Intersection Management 课题，它回答的是高层调度、低层轨迹平滑、安全过滤或真实验证中的一个关键子问题：如何让车辆在路口少停、少冲突、少能耗，同时保持在线可执行。

The paper belongs to the theme of Distributed hierarchical learning for AIM. For a JSSP-based intersection-management thesis, it illuminates one of the key layers: scheduling, smoothing, safety filtering, or real-world validation.

Learning framework lacks explicit optimization/smoothing formulation and formal safety guarantees.

核心科学问题 / Core Scientific Question

如何用层次化长期/短期判别机制改进多智能体交叉口交互策略学习？

How can hierarchical long-term and short-term adversarial discrimination improve multi-agent interaction learning for AIM?

科学贡献 / Scientific Contributions

- 方法贡献 (method contribution): Non-RL adversarial actor/discriminator framework evaluates immediate interaction and final trajectory quality for multi-agent AIM.
- 安全/避碰贡献 (safety contribution): Immediate and final discriminators score interaction safety and trajectory outcomes.
- 平滑/停止避免贡献 (smoothing contribution): Targets reduced travel time and smoother interaction outcomes compared with distributed learning baselines.
- 创新力度评判 (reviewer judgement): 中等偏强：层次化对抗学习有新意，但安全和可解释性仍弱于约束优化方法。

核心概念对照 / Core Concepts

中文概念	English Concept	Role / 学术作用
自动交叉口管理	Autonomous Intersection Management	把信号控制替换为车辆级调度与控制。
轨迹平滑	Trajectory Smoothing	减少急加减速、停车和能耗峰值。
安全约束	Safety Constraints	把碰撞风险写成距离、时间间隔或屏障函数。
Distributed hierarchical adversarial learning	分布式层次化对抗学习	该论文的核心方法族。

教材式学习路径 / Textbook–Style Study Path

- 第一遍先读首页截图、摘要与引言，确认研究问题和场景假设。
- 第二遍沿着技术路线图复现模型变量、约束和求解器输入输出。
- 第三遍逐节阅读 section deep reading，对照 PDF 截图检查公式、图表和实验。
- 第四遍只站在审稿人视角重读实验，判断 baseline、公平性、消融和鲁棒性。
- 最后把博士 checklist 写成自己的复现实验或论文拓展计划。

博士阅读 Checklist / Doctoral Reading Checklist

- 能否独立重写核心公式：层次对抗学习目标 / Hierarchical adversarial–learning objective，并解释每个符号的物理意义？
 - 能否画出输入–模型–算法–控制–验证的数据流图？
 - 能否指出至少两个强假设，并说明它们在真实交通中何时失效？
 - 能否复述实验基准是否公平，指标是否覆盖效率、安全、能耗和计算时间？
 - 能否提出一个与自己 JSSP/RL/smoother 课题直接相连的拓展实验？
-

关键截图讲解 / PDF Snapshot Explanation

截图来自本地 PDF 的精选页面，用于学习定位和版面理解；请与原 PDF 平行阅读，不把截图视作全文替代。





Figure 5: The general framework of hierarchical adversarial learning

实验或结果页 / Experiments or results page

p.9 · figure_crop

这张截图用于把笔记中的抽象概念拉回原论文页面，便于你平行阅读原文、公式、图表和实验设置。

hierarchical adversarial learning to guide the long-term and short-term behaviors of CAVs. The hierarchical adversarial learning isolates immediate interaction discriminator and final trajectory discriminator, where the former arbitrates the priority of each-step interaction and the latter qualifies the entire trajectory. The main contributions of this paper includes:

1. A hierarchical adversarial learning by long-term and short-term discrimination is adopted to realize the complicated interaction among CAVs, which does not depend on design of reward and has low training cost.
2. Potential energy method is applied to balance speed acquisition (efficiency) and conflict-avoidance (safety) in different distances away from the stop-lines.
3. the proposed model is evaluated in a four-way-six-lane intersection and found able to prevent collision and eliminate dead-lock at the crossing zone, and it performs better in safety, efficiency, fairness than the baselines.

The rest of the paper is organized as follows. Section II reviews relevant literature in AIM problem, Section III illustrates details of D-SAL. In Section IV, we conduct experiments to evaluate the performance of D-SAL, and compare it with baselines. Conclusion and discussions are presented in Section V.

2 Literature Review

The framework of Autonomous Intersection Management (AIM) presented by Dresner et al. laid the cornerstone for investigating the interaction of CAVs at intersections [1]. Depending on the decision variables, the AIM can be categorized into priority-based and action-based management. As the name implies, priority-based management produces a crossing order for all approaching vehicles and determines the priority between pairs of conflicting vehicles, while action-based management generates reasonable speed trajectories for CAVs to achieve Pareto optimality among goals like safety, efficiency and fairness. Compared with the action-based approach, the priority-based approach reduces the solution space of the AIM problem, and it can prevent collisions and deadlocks effectively under the restriction of yield-and-go rules. However, action-based approach is more executable in distributed AIM, and the more precise control objectives in action-based way enable more optimal scheduling. In order to combine the advantages of both priority-based and action-based approaches, the hierarchical model employs an upper layer to negotiate priorities and a lower layer to generate optimal velocity trajectories [2]. The model proposed in our study is also a combination of

核心方法或算法页 / Core method or algorithm page

p.2 · full_page

这张截图用于把笔记中的抽象概念拉回原论文页面，便于你平行阅读原文、公式、图表和实验设置。

This snapshot anchors the study note to the original PDF page so you can read the explanation alongside the source layout.

与当前课题的连接：Useful for explaining why learned interaction policies must still be checked by a low-level feasibility/safety layer.

This snapshot anchors the study note to the original PDF page so you can read the explanation alongside the source layout.

与当前课题的连接：Useful for explaining why learned interaction policies must still be checked by a low-level feasibility/safety layer.

In terms of fairness, First-Come-First-Serve seems to have more equitable principle, but it does not perform as well as it could. The method does not make full use of intersection capacity when some conflicting CAVs have adjacent priorities, and some CAVs that can cross simultaneously in the non-conflicting direction are forced to wait before the stop line, leading to large variations in efficiency at different time and with different short-time flows, thus enlarge the travel time differences for CAVs arriving at different instants. The learning-based approach, on the other hand, is incentivized by efficiency terms and reduces forced stops and waits, resulting in more equity in terms of differences in travel time. With regard to fuel consumption, it is mainly caused by two reasons, the stop-and-go behavior of CAVs before crossing the intersection in the signal-based and priority-based management, and the acceleration and deceleration in adjustment of arrival time and crossing speed in the action-based method. Without the inclusion of fuel consumption feedback, the fuel consumption of CAV in the action-based method is more than that in the signal-based method. Among these action-based method, the proposed model has the least consumption due to less speed adjustment in the entire trajectory. In addition, the gap in fuel consumption between the action-based and signal-based methods is shrunk as traffic volume rise.

4.4.3 Interaction Pattern

We find that CAVs tend to pass the crossing zones at unsignalized intersection with some specified patterns, some of which are enumerated in figure 7. Each scenario in figure 7 contains snapshots with a total interval of 3.2s. In Scenario one, CAVs in the north-south, south-south, east-west, west-east interfere through the intersection, and CAVs in other directions intersepte through the gaps between CAVs in those directions. The conflict points between left-turn lanes are more concentrated than those between straight lanes, which are more likely to generate conflicts and deadlocks in unsignalized control. When crossing these conflict points, CAVs tend to adopt strategies like rhythmic control. Each rhythm cycle generally contains three or four CAVs turning left in sequence, as shown in Scenario two to Scenario four in 7. The crossing order among them is usually organized as: CAVs in two non-conflicting directions turn left simultaneously, and then CAVs in the two other non-conflicting directions turn left simultaneously, or four CAVs turn left in a clockwise or counterclockwise direction in sequence. The combined patterns of straight and left turns are not significant at low traffic volumes, but arise when traffic volumes are elevated, i.e., various specified rhythmic patterns of movements arise between CAVs in all lanes, instead of separate crossing patterns between straight and left-turn CAVs respectively. In high traffic flows, platoon formation occurs spontaneously among CAVs in the same direction just like platoons-based method and LQP control, but unlike these two method, the spontaneous formation will enlarge the headway with front CAV to enable the upcoming CAVs in the conflicting direction to pass as shown in Scenario Five7, reducing the waiting time and unnecessary slowdown and stop of conflicting CAVs, and further improving the efficiency of passage.

5 Conclusion

In this study, we develop a non-RL learning-based framework for addressing AIM problem with pure CAVs. The framework applies hierarchical adversarial learning to regulate the instant and farighted behaviors of CAVs with immediate discriminators and a final discriminator. To the best of our knowledge, this study is one of the first attempts to decouple the temporal dimension from learning-based methods, and enable CAVs to concentrate more on the immediate interaction and the final results of the interaction, which is found to be robust to conflict avoidance. To balance safety and efficiency, a dynamic potential energy term is used to weigh the importance of efficiency at different distances from the crossing zone and on different stages of training process. Besides, two important mechanisms – state maintenance and action mask enable the actor network to find the optimal policy more stably in the interaction.

The experiments to evaluate and validate the proposed model are configured at a four-way-six-lane intersections, the results manifest that the proposed model: (1) has satisfactory convergence on the losses of actor network and discriminator networks. (2) can navigate CAVs to intersepte and interact efficiently during intersection passage. Besides, several methods commonly used for intersection scheduling management are compared with the proposed model, the comparisons show that the proposed model can achieve the fastest passage while ensuring safety.

However, since the action-based method involves frequent acceleration and deceleration, fuel consumption and comfort may inferior to some priority-based and signal-based methods, which need to be further enhanced in the future research, possibly by involving appropriate relevant objectives in the training.

References

[1] Kurt Dresner and Peter Stone. Multiagent traffic management: A reservation-based intersection control mechanism. In *Autonomous Agents and Multiagent Systems, International Joint Conference on*, volume 3, pages 530-537.

结论或局限页 / Conclusion or limitations page

p.14 · full_page

这张截图用于把笔记中的抽象概念拉回原论文页面，便于你平行阅读原文、公式、图表和实验设置。

This snapshot anchors the study note to the original PDF page so you can read the explanation alongside the source layout.

与当前课题的连接：Useful for explaining why learned interaction policies must still be checked by a low-level feasibility/safety layer.

章节精读 / Section-by-Section Deep Reading

p.1 · motivation

1 Introduction

章节目标：建立研究动机：为什么传统信号控制、FCFS 或单车控制不足以处理该论文关注的交叉口协同问题。

Learning goal: Understand why the paper needs a new coordination method rather than a standard signal, FCFS, or single-vehicle controller.

关键论点 / Key argument: 这一部分通常把痛点落到 Distributed hierarchical learning for AIM：既要提高效率，又要保持安全与在线可执行。

技术焦点 / Technical focus: 精读时标出作者批判的 baseline、场景假设、交通参与者类型和隐含优化目标。

审稿追问 / Reviewer question: 作者指出的痛点是否真的需要新方法，还是可以由更强的优化器/更保守规则解决？

p.2 · bridge

1. A hierarchical adversarial learning by long-term and short-term discrimination is adopted to realize the

章节目标：承接前后章节，补充局部定义、案例或实现细节。

Learning goal: Recover implementation details needed for reproduction.

关键论点 / Key argument: 这类章节常常不是主贡献，但决定你能否真正复现方法。

技术焦点 / Technical focus: 检查它是否给出了参数、伪代码、场景划分或实现细节。

审稿追问 / Reviewer question: 跳过该节会不会导致公式或实验无法复现？

p.2 · method

2. Potential energy method is applied to balance speed acquisition (efficiency) and conflict-avoidance (safety) in

章节目标：解释算法如何把模型转化为可执行决策，并说明在线运行机制。

Learning goal: Trace how the paper turns the model into executable online decisions.

关键论点 / Key argument: D-HAL 通过长期和短期判别器引导策略学习，让智能体既模仿/区分整体交互轨迹，又修正局部短时动作。层次化目标试图改善多智能体交互中的稀疏奖励和非平稳性。

技术焦点 / Technical focus: 画出输入、核心求解器、输出和安全检查接口，明确哪些步骤是优化、哪些是规则、哪些是学习。

审稿追问 / Reviewer question: 若算法某一步失败，论文有没有 fallback、重规划或安全过滤机制？

p.2 · context

2 Literature Review

章节目标：建立方法谱系和相关工作边界，避免把本文贡献误读为完全从零开始。

Learning goal: Place the paper inside the broader method taxonomy.

关键论点 / Key argument: 该部分说明作者站在哪些已有工作之上，以及本文选择绕开或继承了哪些限制。

技术焦点 / Technical focus: 把相关工作按 scheduler、smoother、safety filter、evaluation 四类重排。

审稿追问 / Reviewer question: 作者是否公平比较了最接近的强 baseline？

3 Methodology

章节目标：解释算法如何把模型转化为可执行决策，并说明在线运行机制。

Learning goal: Trace how the paper turns the model into executable online decisions.

关键论点 / Key argument: D-HAL 通过长期和短期判别器引导策略学习，让智能体既模仿/区分整体交互轨迹，又修正局部短时动作。层次化目标试图改善多智能体交互中的稀疏奖励和非平稳性。

技术焦点 / Technical focus: 画出输入、核心求解器、输出和安全检查接口，明确哪些步骤是优化、哪些是规则、哪些是学习。

审稿追问 / Reviewer question: 若算法某一步失败，论文有没有 fallback、重规划或安全过滤机制？

3.1 Problem Statement

章节目标：把自然语言交通问题压缩为变量、状态、约束、目标函数或图结构。

Learning goal: Translate the traffic problem into variables, states, constraints, objectives, or graph relations.

关键论点 / Key argument: 本节是复现论文的核心入口，应与公式“层次对抗学习目标 / Hierarchical adversarial-learning objective”一起读。

技术焦点 / Technical focus: 逐项检查车辆状态、冲突关系、控制输入、时间/空间索引和安全阈值是否定义完整。

审稿追问 / Reviewer question: 这些建模假设在混合交通、通信延迟、感知误差和高密度车流下是否仍成立？

3.3 Reservation Mechanism

章节目标：承接前后章节，补充局部定义、案例或实现细节。

Learning goal: Recover implementation details needed for reproduction.

关键论点 / Key argument: 这类章节常常不是主贡献，但决定你能否真正复现方法。

技术焦点 / Technical focus: 检查它是否给出了参数、伪代码、场景划分或实现细节。

审稿追问 / Reviewer question: 跳过该节会不会导致公式或实验无法复现？

3.4 Hierarchical Adversarial Learning

章节目标：承接前后章节，补充局部定义、案例或实现细节。

Learning goal: Recover implementation details needed for reproduction.

关键论点 / Key argument: 这类章节常常不是主贡献，但决定你能否真正复现方法。

技术焦点 / Technical focus: 检查它是否给出了参数、伪代码、场景划分或实现细节。

审稿追问 / Reviewer question: 跳过该节会不会导致公式或实验无法复现？

逐段式导读 / Paragraph-Level Walkthrough

Opening motivation / 开篇动机段

先把作者的痛点翻译成一句博士问题：如何用层次化长期/短期判别机制改进多智能体交叉口交互策略学习？读这一段时不要急着接受动机，要追问传统 FCFS、信号灯或保守 MPC 为什么不足。

Turn the opening motivation into one doctoral-level question and test whether the paper truly needs a new method.

Assumption and setting / 假设与场景段

把车辆类型、通信条件、路口类型和动力学假设列成表。该文的关键边界包括：专家轨迹或高质量行为分布可用于判别学习。；长期/短期判别目标与真实交通安全效率一致。；仿真环境足以覆盖多智能体交互复杂性。

List vehicle type, communication, intersection setting, and dynamics assumptions before reading the equations.

Model and notation / 模型与符号段

围绕核心公式“层次对抗学习目标 / Hierarchical adversarial-learning objective”复写符号表，确认每个变量是决策变量、状态变量、参数还是约束阈值。

Rewrite the symbol table around the core formulation and classify variables as states, decisions, parameters, or constraints.

Algorithm mechanism / 算法机制段

把算法读成一条数据流：输入交通状态，执行 Non-RL adversarial actor/discriminator framework evaluates immediate interaction and final trajectory quality for multi-agent AIM., 输出可执行通行/控制决策，并检查安全接口。

Read the algorithm as a dataflow from traffic state to executable sequence/control decision, with explicit safety interfaces.

Experiment and evidence / 实验证据段

不要只看作者报告的提升，要检查基准、公平性和指标。本文验证描述为：Four-way six-lane intersection experiments against state-of-the-art methods.

Go beyond reported gains and inspect baselines, fairness, metrics, and computation time.

Reviewer synthesis / 审稿式总结段

最后把贡献拆成可发表与需补强两部分。对你课题最直接的用途是：Useful for explaining why learned interaction policies must still be checked by a low-level feasibility/safety layer.

Separate publishable contribution from missing evidence, then connect the result to your own thesis architecture.

技术路线图与贡献量评估 / Technical Route & Contribution Matrix

1. 输入 Input 车辆状态、路径/车道、交通需求、信号或协同信息。	2. 建模 Modeling 层次对抗学习目标 / Hierarchical adversarial-learning objective	3. 核心求解 Core solver Non-RL adversarial actor/discriminator framework evaluates immediate interaction and final trajectory quality for multi-agent AIM.
4. 安全接口 Safety interface Immediate and final discriminators score interaction safety and trajectory outcomes.	5. 平滑与效率 Smoothing and efficiency Targets reduced travel time and smoother interaction outcomes compared with distributed learning baselines.	6. 验证 Validation Four-way six-lane intersection experiments against state-of-the-art methods.

贡献量评估 / Contribution Strength Matrix

Dimension / 维度	Score / 评分	Comment / 评语
理论贡献 / Theoretical contribution	3/5	层次对抗学习目标 / Hierarchical adversarial-learning objective
算法贡献 / Algorithmic contribution	5/5	Non-RL adversarial actor/discriminator framework evaluates immediate interaction and final trajectory quality for multi-agent AIM.
实验贡献 / Experimental contribution	3/5	Four-way six-lane intersection experiments against state-of-the-art methods.
工程贡献 / Engineering contribution	3/5	Immediate and final discriminators score interaction safety and trajectory outcomes.
博士课题价值 / PhD relevance	4/5	Useful for explaining why learned interaction policies must still be checked by a low-level feasibility/safety layer.

2. 理论方法论深度解构 / Deep Dive into Methodology & Theoretical Framework

问题数学建模 / Mathematical Formulation

层次对抗学习目标 / Hierarchical adversarial-learning objective

$$\min_{\pi} \max_{\{D_L, D_S\}} \mathbb{E}_{\tau \sim \pi} [\log D(\tau)] + \mathbb{E}_{\tau \sim \pi} [\log(1 - D(\tau))] - \lambda J_{RL}(\pi)$$

Symbol / 符号	含义	Meaning	Role / 建模作用
π	多智能体策略	multi-agent policy	学习车辆交互行为
D_L	长期判别器	long-term discriminator	评价轨迹级行为质量
D_S	短期判别器	short-term discriminator	评价局部交互动作
τ	轨迹片段	trajectory segment	对抗学习样本
J_{RL}	强化学习收益	RL return	保留任务奖励优化目标

核心算法架构 / Core Algorithm Layout

D-HAL 通过长期和短期判别器引导策略学习，让智能体既模仿/区分整体交互轨迹，又修正局部短时动作。层次化目标试图改善多智能体交互中的稀疏奖励和非平稳性。

D-HAL uses long-term and short-term discriminators to shape multi-agent policies, improving interaction learning under sparse rewards and non-stationarity.

收敛性、稳定性或实时性 / Convergence, Stability, Real-Time Feasibility

对抗训练可能提升样本效率，但也会带来训练不稳定；安全仍主要通过奖励或环境终止惩罚表达，缺少硬约束。

Adversarial training can improve sample efficiency but may be unstable; safety is still mostly reward-based unless paired with constraints.

潜在假设与边界条件 / Assumptions & Boundary Conditions

- 专家轨迹或高质量行为分布可用于判别学习。
- 长期/短期判别目标与真实交通安全效率一致。
- 仿真环境足以覆盖多智能体交互复杂性。

3. 实验验证与结果评判 / Experimental Validation & Result Evaluation

基准对比与环境 / Baselines & Environment

普通 MARL、模仿学习、无层次判别器版本、无对抗损失版本。

Four-way six-lane intersection experiments against state-of-the-art methods.

核心指标 / Metrics

碰撞率、平均延误、通行效率、成功率、奖励曲线和训练样本效率。

消融实验 / Ablation

必须拆分长期判别器、短期判别器、对抗损失和 RL 任务损失；否则难以证明层次结构必要。

博士阅读提醒 / Reading Reminder

阅读实验时不要只看平均值；应检查交通密度、车流方向、随机种子、失败案例和计算耗时。For doctoral reading, inspect density regimes, demand patterns, random seeds, failure cases, and computation time rather than only mean improvements.

图表阅读线索 / Figure and Table Reading Cues

PDF extraction detected approximately 8 figure references and 6 table references. Exact captions remain in the original PDF; the bilingual cues below tell you what to inspect first.

图表名称	Figure/Table Name	Reading focus / 阅读重点
系统框架图	system framework diagram	先确认输入、决策层、控制层和安全约束之间的数据流。
策略网络/训练流程图	policy-network or training-pipeline diagram	检查状态、动作、奖励、经验回放和安全惩罚是否闭环一致。
实验指标对比图表	experimental metric comparison plots/tables	重点看平均值之外的密度变化、失败场景、计算时间和方差。

4. 审稿人视角：批判性思考与未来方向 / Critical Thinking & Future Directions

论文局限性 / Limitations & Weaknesses

- 对抗训练难调参，复现实验成本高。
- 策略可解释性弱，审稿人会追问安全证书。
- 分布式学习结果可能依赖仿真器和奖励设置。

博士课题拓展点 / Research Extensions for PhD

- 用 JSSP 求解器生成专家排序轨迹，训练 imitation/adversarial scheduler。
- 把判别器输出作为调度质量估计，而不是直接控制车辆。
- 在对抗学习外加 CBF/MPC 安全约束，分离效率学习与安全保证。

如何服务当前课题 / Use in This Project

Useful for explaining why learned interaction policies must still be checked by a low-level feasibility/safety layer.

5. 交互式可视化与 PDF 排版蓝图 / Interactive Visualization & PDF Layout Blueprint

动态参数探索 / Parameter Exploration

对抗损失权重 / Adversarial loss weight:

权重适中可提升交互拟合，过高会牺牲任务奖励稳定性。

Moderate weight improves interaction shaping; too much can destabilize RL.

HTML 结构化布局建议 / HTML Structuring

- 顶部固定 metadata bar：题名、作者、年份、主题、PDF 页数。
- 左侧目录/section navigator，右侧为五大精读模块。
- 公式卡片使用 `<pre>` 或 LaTeX block，紧跟符号表。
- 审稿人批注用 reviewer-note block，博士拓展用 research-idea list。
- 交互参数用 `input[type=range]` + canvas 曲线，打印 PDF 时隐藏交互控件但保留解释。

来源锚点与逐节阅读路径 / Source Anchors & Section-by-Section Reading Map

以下锚点来自本地 PDF 抽取，用于回到原文复核；笔记不保存论文全文。

Page	Extracted heading
p.1	1 Introduction
p.2	2. Potential energy method is applied to balance speed acquisition (efficiency) and conflict-avoidance (safety) in
p.9	4 Experiments

p.1 · 1 Introduction

定位问题背景、研究缺口和为什么现有保守/非协同方法不足。
Frames the motivation, research gap, and limits of conservative or non-cooperative methods.

p.2 · 1. A hierarchical adversarial learning by long-term and short-term discrimination is adopted to realize the

作为局部论证段落阅读，关注它如何服务主问题。
Read as a local argument and ask how it supports the main question.

p.2 · 2. Potential energy method is applied to balance speed acquisition (efficiency) and conflict-avoidance (safety) in

给出算法流程和求解机制，应重点追踪输入、输出和安全接口。
Gives the algorithmic mechanism; track inputs, outputs, and safety interfaces.

p.2 · 2 Literature Review

作为局部论证段落阅读，关注它如何服务主问题。
Read as a local argument and ask how it supports the main question.

p.4 · 3 Methodology

给出算法流程和求解机制，应重点追踪输入、输出和安全接口。
Gives the algorithmic mechanism; track inputs, outputs, and safety interfaces.

p.4 · 3.1 Problem Statement

把交通交互转写为状态、约束、目标或图结构，是精读时最应复现的部分。
Defines states, constraints, objectives, or graph structures; this is the section to reproduce carefully.

p.5 · 3.3 Reservation Mechanism

作为局部论证段落阅读，关注它如何服务主问题。

Read as a local argument and ask how it supports the main question.

p.6 · 3.4 Hierarchical Adversarial Learning

作为局部论证段落阅读，关注它如何服务主问题。

Read as a local argument and ask how it supports the main question.
